

Phase 1

Data Cleaning

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.simplefilter("ignore")
```


```
In [2]: df =pd.read_csv("AviationData.csv",encoding = "latin1")
```

```
In [3]: df.head()
```

Out[3]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States

5 rows × 31 columns




```
In [4]: df.tail()
```

```
Out[4]:
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country
88884	20221227106491	Accident	ERA23LA093	2022-12-26	Annapolis, MD	United States
88885	20221227106494	Accident	ERA23LA095	2022-12-26	Hampton, NH	United States
88886	20221227106497	Accident	WPR23LA075	2022-12-26	Payson, AZ	United States
88887	20221227106498	Accident	WPR23LA076	2022-12-26	Morgan, UT	United States
88888	20221230106513	Accident	ERA23LA097	2022-12-29	Athens, GA	United States

5 rows × 31 columns




```
In [5]: df.sample(5)
```

```
Out[5]:
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country
56400	20040331X00405	Accident	LAX04LA170	2004-03-25	Planada, CA	United States
40906	20001208X07000	Accident	CHI97FA030	1996-11-18	GRAND RAPIDS, MI	United States
84170	20191107X55010	Accident	WPR20FA019	2019-11-07	Upland, CA	United States
27909	20001212X16974	Accident	DEN91FA067	1991-05-06	HARTLEY, TX	United States
57183	20040830X01320	Accident	CHI04CA210	2004-07-31	Plymouth, MI	United States

5 rows × 31 columns



In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                    88889 non-null  object
2   Accident.Number                      88889 non-null  object
3   Event.Date                           88889 non-null  object
4   Location                             88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                             34382 non-null  object
7   Longitude                            34373 non-null  object
8   Airport.Code                         50132 non-null  object
9   Airport.Name                         52704 non-null  object
10  Injury.Severity                      87889 non-null  object
11  Aircraft.damage                      85695 non-null  object
12  Aircraft.Category                    32287 non-null  object
13  Registration.Number                 87507 non-null  object
14  Make                                88826 non-null  object
15  Model                               88797 non-null  object
16  Amateur.Built                       88787 non-null  object
17  Number.of.Engines                   82805 non-null  float64
18  Engine.Type                         81793 non-null  object
19  FAR.Description                     32023 non-null  object
20  Schedule                            12582 non-null  object
21  Purpose.of.flight                   82697 non-null  object
22  Air.carrier                         16648 non-null  object
23  Total.Fatal.Injuries                77488 non-null  float64
24  Total.Serious.Injuries              76379 non-null  float64
25  Total.Minor.Injuries                76956 non-null  float64
26  Total.Uninjured                     82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight               61724 non-null  object
29  Report.Status                       82505 non-null  object
30  Publication.Date                    75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

In [7]: df.describe()

Out[7]:


	Number.of.Engines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries	Total.I
count	82805.000000	77488.000000	76379.000000	76956.000000	8295
mean	1.146585	0.647855	0.279881	0.357061	
std	0.446510	5.485960	1.544084	2.235625	2
min	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	0.000000	0.000000	0.000000	
50%	1.000000	0.000000	0.000000	0.000000	
75%	1.000000	0.000000	0.000000	0.000000	
max	8.000000	349.000000	161.000000	380.000000	65

In [8]: `df.describe(include = "O")`

Out[8]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Coi
count	88889	88889	88889	88889	88837	8
unique	87951	2	88863	14782	27758	
top	20001212X19172	Accident	CEN22LA149	1984-06-30	ANCHORAGE, AK	U S
freq	3	85015	2	25	434	8

4 rows × 26 columns



In [9]: `#column name`
`df.columns = df.columns.str.lower().str.replace(r'\.', '_', regex=True)`

In [10]: `df.columns`

Out[10]: Index(['event_id', 'investigation_type', 'accident_number', 'event_date',
'location', 'country', 'latitude', 'longitude', 'airport_code',
'airport_name', 'injury_severity', 'aircraft_damage',
'aircraft_category', 'registration_number', 'make', 'model',
'amateur_built', 'number_of_engines', 'engine_type', 'far_descripti
on',
'schedule', 'purpose_of_flight', 'air_carrier', 'total_fatal_injuri
es',
'total_serious_injuries', 'total_minor_injuries', 'total_uninjure
d',
'weather_condition', 'broad_phase_of_flight', 'report_status',
'publication_date'],
dtype='object')

```
In [11]: df.dtypes
```

```
Out[11]: event_id                object
investigation_type             object
accident_number                object
event_date                     object
location                       object
country                        object
latitude                       object
longitude                      object
airport_code                   object
airport_name                   object
injury_severity                object
aircraft_damage                object
aircraft_category              object
registration_number            object
make                           object
model                           object
amateur_built                  object
number_of_engines              float64
engine_type                     object
far_description                 object
schedule                       object
purpose_of_flight              object
air_carrier                    object
total_fatal_injuries            float64
total_serious_injuries          float64
total_minor_injuries            float64
total_uninjured                 float64
weather_condition               object
broad_phase_of_flight           object
report_status                   object
publication_date                object
dtype: object
```

Checking Missing Values and Correcting It.



```
In [12]: df.isnull().sum()
```

```
Out[12]: event_id                0
investigation_type              0
accident_number                 0
event_date                     0
location                       52
country                       226
latitude                      54507
longitude                     54516
airport_code                   38757
airport_name                   36185
injury_severity                1000
aircraft_damage                3194
aircraft_category              56602
registration_number            1382
make                           63
model                          92
amateur_built                  102
number_of_engines              6084
engine_type                    7096
far_description                56866
schedule                       76307
purpose_of_flight              6192
air_carrier                    72241
total_fatal_injuries           11401
total_serious_injuries         12510
total_minor_injuries           11933
total_uninjured                5912
weather_condition              4492
broad_phase_of_flight          27165
report_status                  6384
publication_date               13771
dtype: int64
```

```
In [13]: df.columns
```

```
Out[13]: Index(['event_id', 'investigation_type', 'accident_number', 'event_date',
               'location', 'country', 'latitude', 'longitude', 'airport_code',
               'airport_name', 'injury_severity', 'aircraft_damage',
               'aircraft_category', 'registration_number', 'make', 'model',
               'amateur_built', 'number_of_engines', 'engine_type', 'far_descripti
on',
               'schedule', 'purpose_of_flight', 'air_carrier', 'total_fatal_injuri
es',
               'total_serious_injuries', 'total_minor_injuries', 'total_uninjure
d',
               'weather_condition', 'broad_phase_of_flight', 'report_status',
               'publication_date'],
              dtype='object')
```

```
In [14]: df = df[['event_id', 'event_date',  
                'location', 'country', 'airport_code',  
                'airport_name', 'injury_severity', 'aircraft_damage',  
                'aircraft_category', 'make', 'model',  
                'amateur_built', 'number_of_engines', 'engine_type',  
                'purpose_of_flight', 'air_carrier', 'total_fatal_injuries',  
                'total_serious_injuries', 'total_minor_injuries', 'total_uninjured',  
                'weather_condition', 'broad_phase_of_flight',  
                ]]
```

```
In [15]: df.columns
```

```
Out[15]: Index(['event_id', 'event_date', 'location', 'country', 'airport_code',  
               'airport_name', 'injury_severity', 'aircraft_damage',  
               'aircraft_category', 'make', 'model', 'amateur_built',  
               'number_of_engines', 'engine_type', 'purpose_of_flight', 'air_carri  
er',  
               'total_fatal_injuries', 'total_serious_injuries',  
               'total_minor_injuries', 'total_uninjured', 'weather_condition',  
               'broad_phase_of_flight'],  
              dtype='object')
```

```
In [16]: # Forward fill, then backward fill  
df['event_date'] = df['event_date'].fillna(method='ffill').fillna(method='b
```

```
In [17]: # Impute using mode (for categorical data)  
for col in df.select_dtypes(include = ["O"]):  
    mode_value = df[col].mode()[0]  
    df[col].fillna(mode_value, inplace=True)
```

```
In [18]: for col in df.select_dtypes(include = ["number"]):  
    mean_value = df[col].mean()  
    df[col].fillna(mean_value, inplace=True)
```

```
In [19]: df.isnull().sum()
```

```
Out[19]: event_id          0
         event_date       0
         location         0
         country          0
         airport_code     0
         airport_name     0
         injury_severity  0
         aircraft_damage  0
         aircraft_category 0
         make             0
         model            0
         amateur_built    0
         number_of_engines 0
         engine_type      0
         purpose_of_flight 0
         air_carrier       0
         total_fatal_injuries 0
         total_serious_injuries 0
         total_minor_injuries 0
         total_uninjured   0
         weather_condition 0
         broad_phase_of_flight 0
         dtype: int64
```

```
In [20]: for col in df.select_dtypes(include = ["O"]):df[col] = df[col].str.lower()
```

```
In [21]: df.duplicated().sum()
```

```
Out[21]: 25
```

```
In [22]: df.drop_duplicates(inplace=True)
```

```
In [23]: df.duplicated().sum()
```

```
Out[23]: 0
```



```
In [24]: numeric_df = df.select_dtypes(include=[np.number])
numeric_df
```

Out[24]:

	number_of_engines	total_fatal_injuries	total_serious_injuries	total_minor_injuries	total_injuries
0	1.000000	2.0	0.000000	0.000000	2.0
1	1.000000	4.0	0.000000	0.000000	4.0
2	1.000000	3.0	0.279881	0.357061	3.636942
3	1.000000	2.0	0.000000	0.000000	2.0
4	1.146585	1.0	2.000000	0.357061	4.503646
...
88884	1.146585	0.0	1.000000	0.000000	1.000000
88885	1.146585	0.0	0.000000	0.000000	0.000000
88886	1.000000	0.0	0.000000	0.000000	0.000000
88887	1.146585	0.0	0.000000	0.000000	0.000000
88888	1.146585	0.0	1.000000	0.000000	1.000000

88864 rows × 5 columns

Checking for Outliers

```

In [25]: # Grid layout
rows, cols = 2, 3
fig, axes = plt.subplots(rows, cols, figsize=(21, 13))

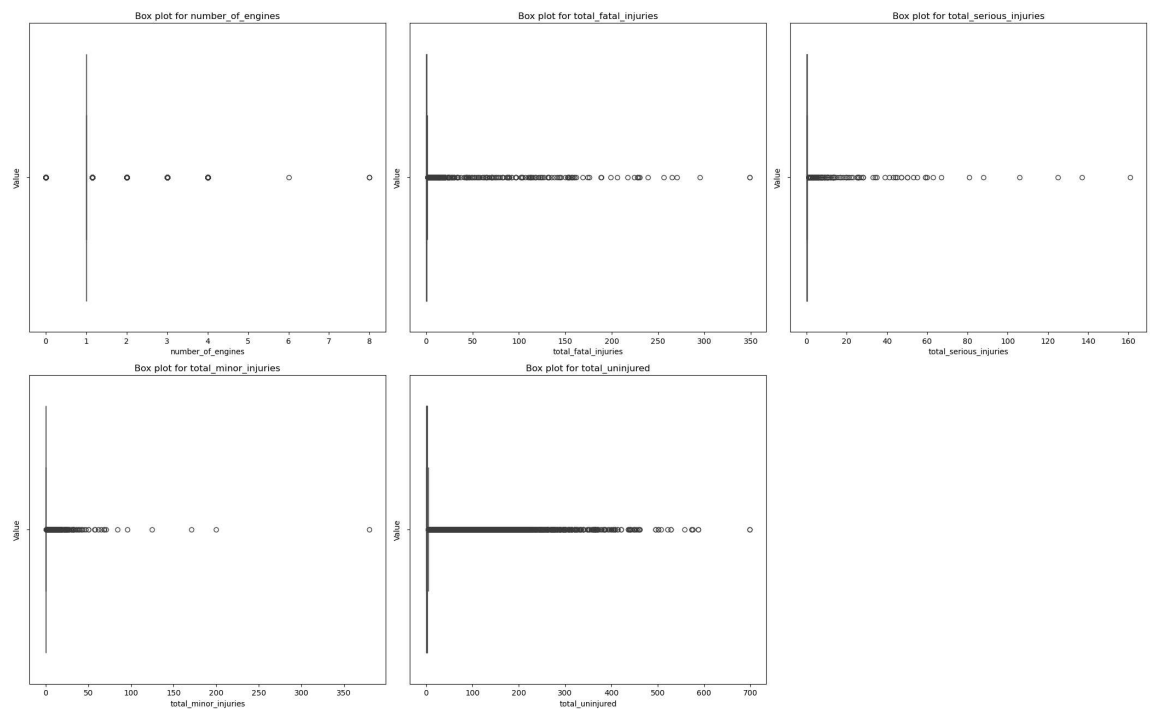
# Flatten
axes = axes.flatten()

for i, column in enumerate(numeric_df):
    sns.boxplot(x=df[column], ax = axes[i])
    axes[i].set_title(f"Box plot for {column}")
    axes[i].set_xlabel(column)
    axes[i].set_ylabel('Value')

# Hide empty subplots
for j in range(i + 1, rows * cols):
    axes[j].axis('off')

plt.tight_layout()
plt.show()

```



```
In [26]: # Select only numeric columns
numeric_data = df.select_dtypes(include=['number'])

# Calculate q1, q3, and IQR for numeric data only
q1 = numeric_data.quantile(0.25) # .25
q3 = numeric_data.quantile(0.75) # .75
IQR = q3 - q1

# Lower and upper bounds for identifying outliers
lower_bound = q1 - (1.5 * IQR)
upper_bound = q3 + (1.5 * IQR)

# Identify outliers in numeric columns
outliers_ = set()
for col in numeric_data.columns:
    outliers = numeric_data[(numeric_data[col] < lower_bound[col]) | (numeric_data[col] > upper_bound[col])]
    outliers_.update(outliers.index)

# Count number of rows before removing outliers
num_rows_before = len(df)

# Remove rows with outliers from the original DataFrame (important!)
df = df.drop(index=outliers_)

# Num of rows after removing outliers
num_rows_after = len(df)

# Number of rows removed
rows_removed = num_rows_before - num_rows_after

# Print the results
print(f"Number of rows before removing outliers: {num_rows_before}")
print(f"Number of rows after removing outliers: {num_rows_after}")
print(f"Number of rows removed: {rows_removed}")
```

Number of rows before removing outliers: 88864

Number of rows after removing outliers: 42604

Number of rows removed: 46260

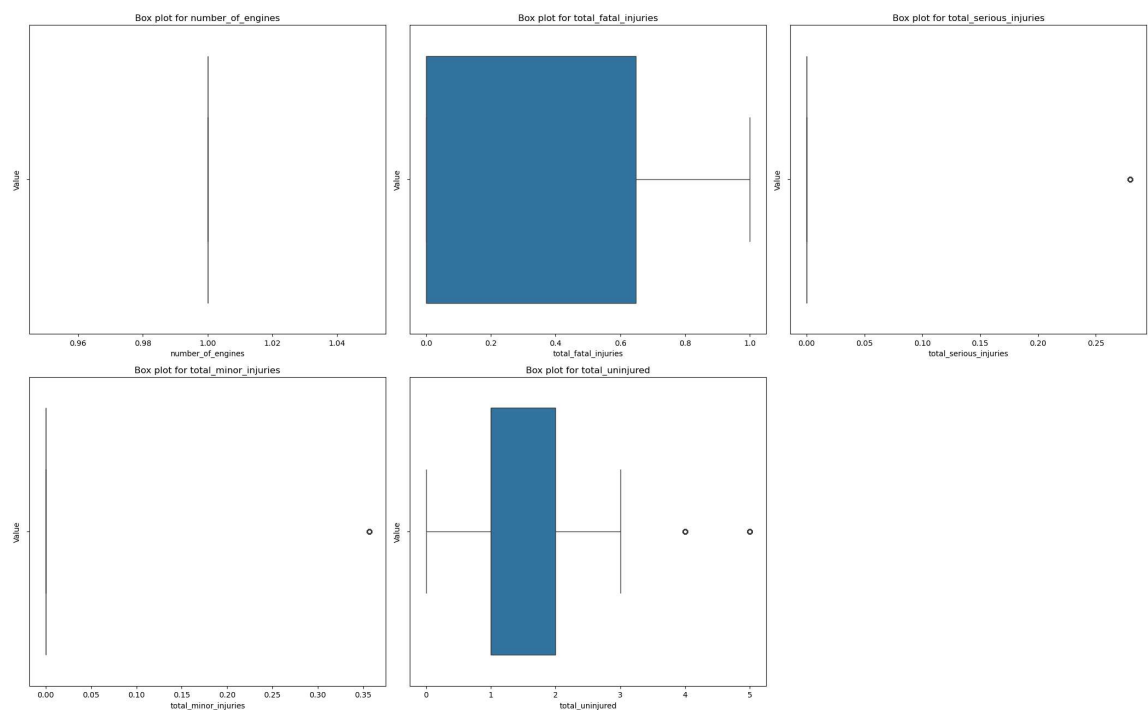
```
In [27]: rows, cols = 2, 3
fig, axes = plt.subplots(rows, cols, figsize=(21, 13))

# Flatten
axes = axes.flatten()

for i, column in enumerate(numeric_df):
    sns.boxplot(x=df[column], ax = axes[i])
    axes[i].set_title(f"Box plot for {column}")
    axes[i].set_xlabel(column)
    axes[i].set_ylabel('Value')

# Hide empty subplots
for j in range(i + 1, rows * cols):
    axes[j].axis('off')

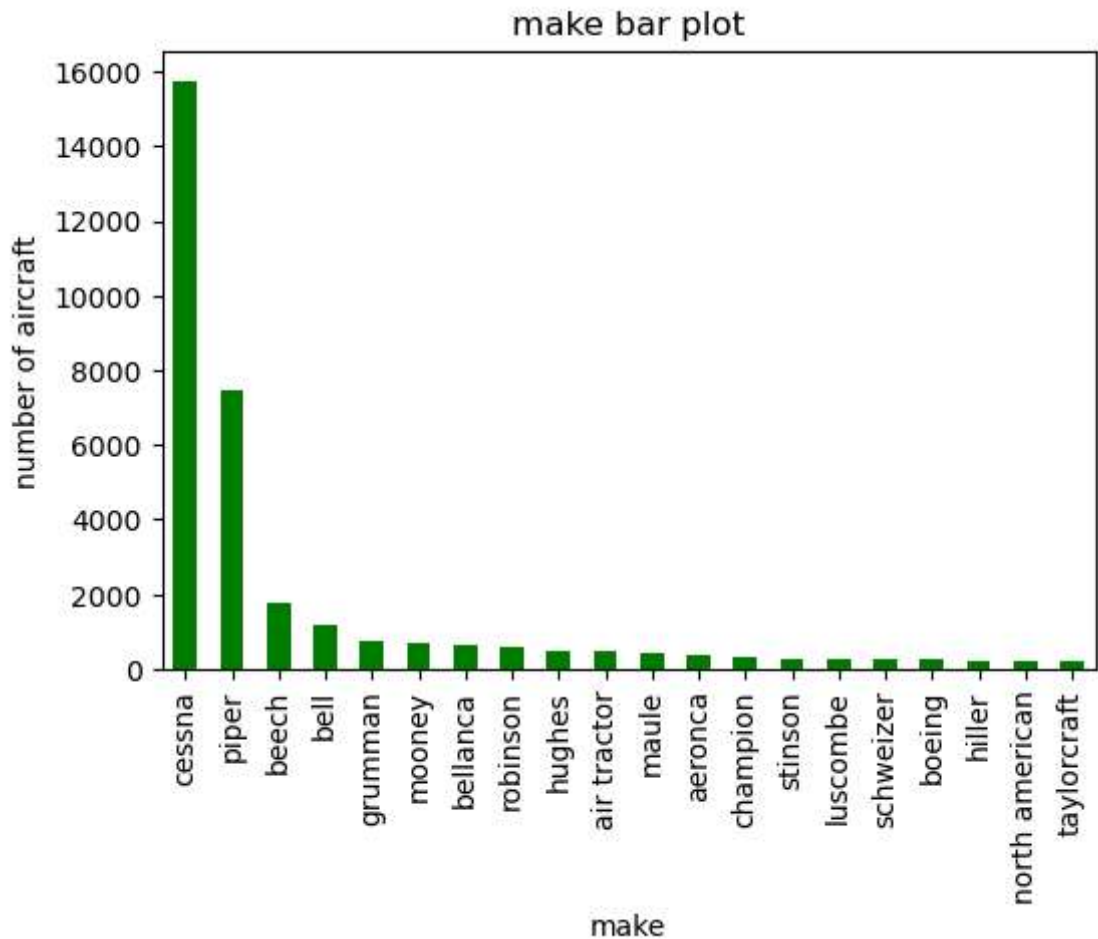
plt.tight_layout()
plt.show()
```




Saving the cleaned data.

```
In [28]: df.to_csv("Project_1R.csv", index = False)
```

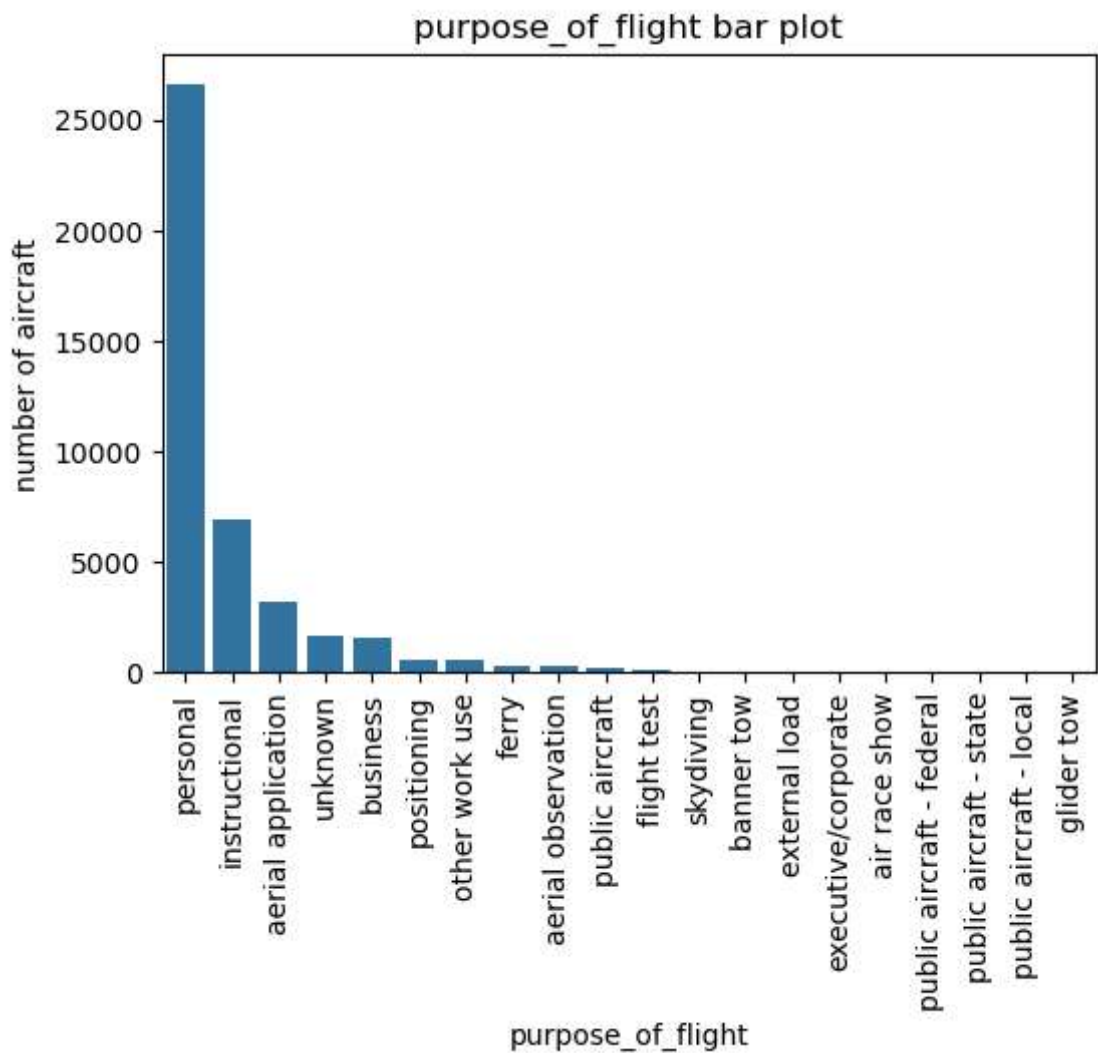
```
In [31]: make_count = df["make"].value_counts().iloc[:20]
make_count
#visual using matplotlib bar
plt.figure(figsize=(6,4))
make_count.plot(kind="bar", color='g')
plt.title("make bar plot")
plt.ylabel("number of aircraft")
plt.show()
```



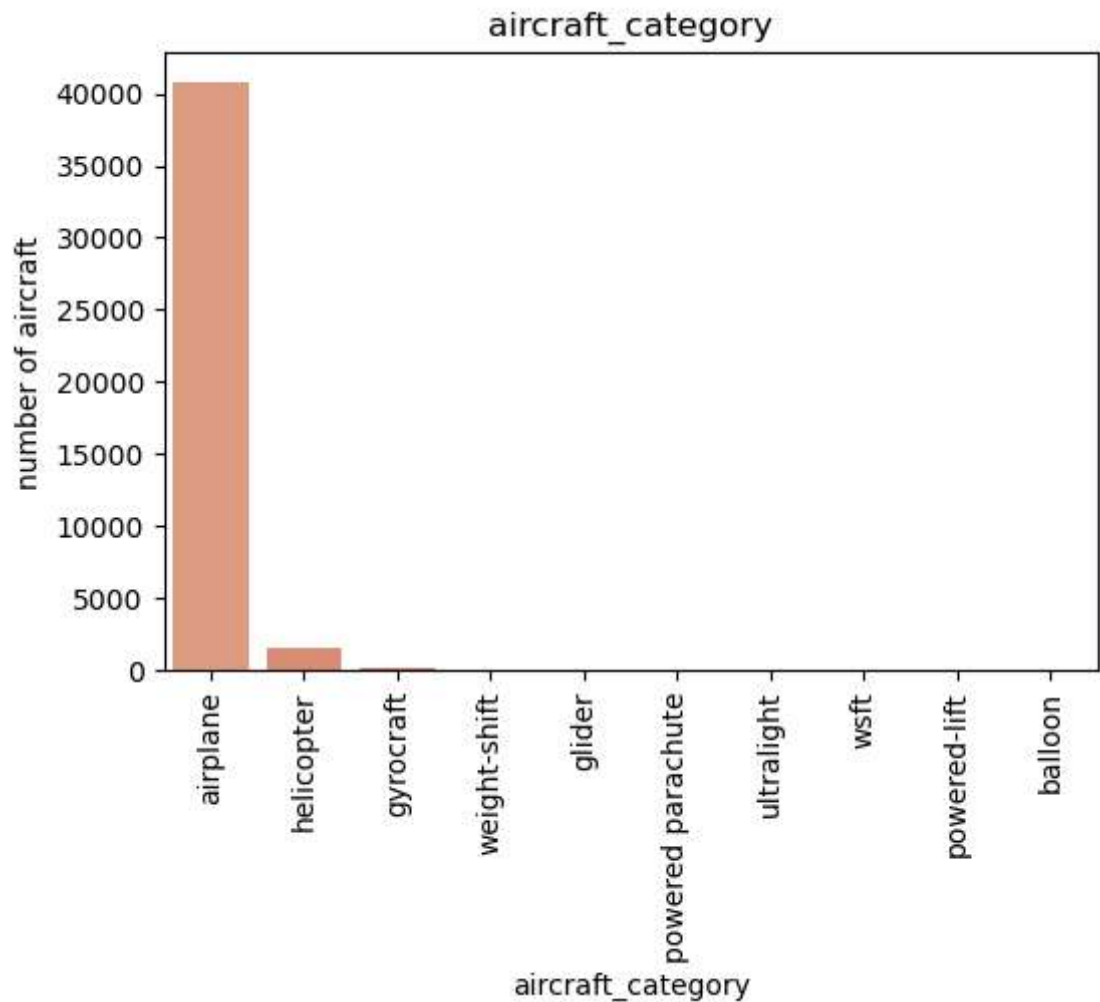
```
In [32]: model_count = df["model"].value_counts().iloc[:20]
model_count
#visual using matplotlib bar
plt.figure(figsize=(6,4))
model_count.plot(kind="bar", color='g')
plt.title("model bar plot")
plt.ylabel("number of aircraft")
plt.show()
```



```
In [38]: purpose_of_flight_count = df["purpose_of_flight"].value_counts().reset_index()
purpose_of_flight_count
#visual using matplotlib bar
plt.figure(figsize=(6,4))
sns.barplot(x = purpose_of_flight_count["purpose_of_flight"],y = purpose_of_flight_count["count"])
plt.title("purpose_of_flight bar plot")
plt.ylabel("number of aircraft")
plt.xticks(rotation = 90)
plt.show()
```

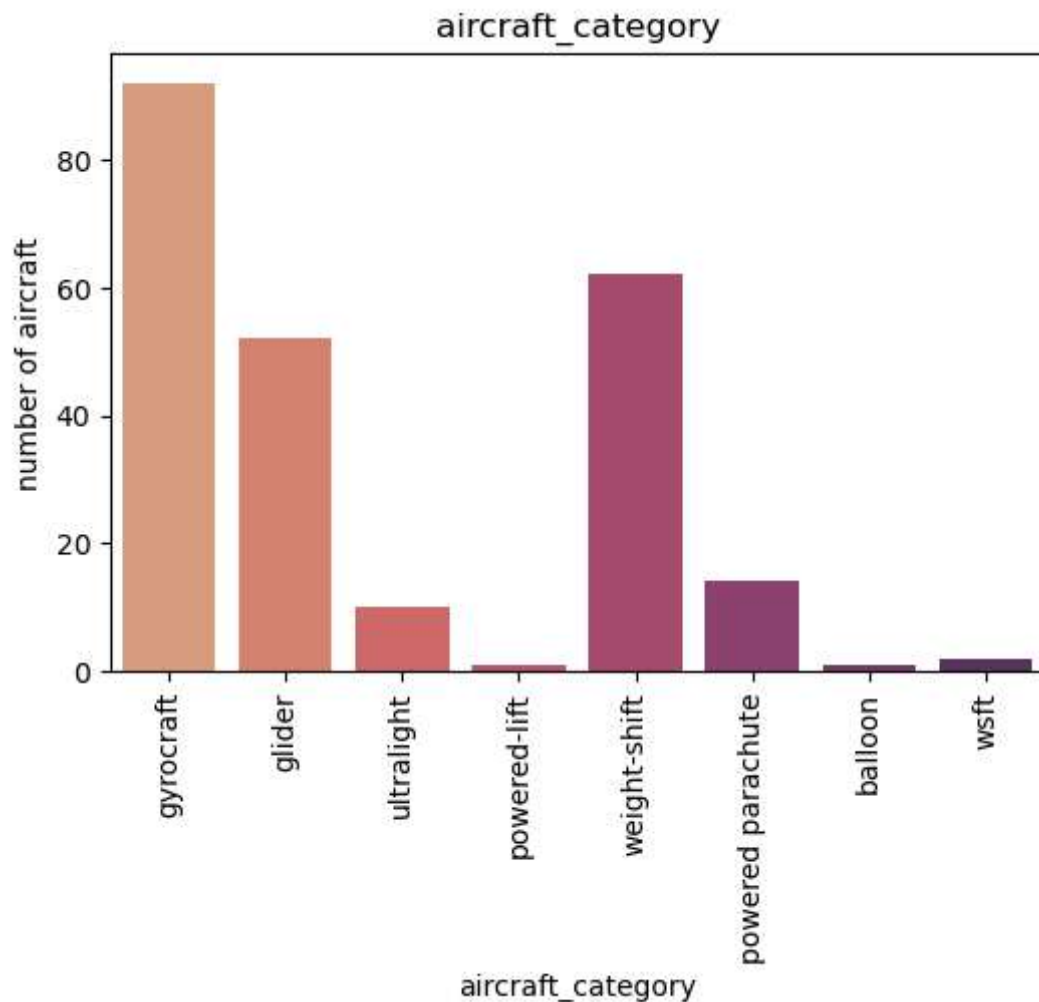


```
In [59]: df = df[df['aircraft_category'] != 'unknown']
aircraft_category = df["aircraft_category"].value_counts().reset_index().il
aircraft_category
#visual using matplotlib bar
plt.figure(figsize=(6,4))
sns.barplot(data = aircraft_category, x = "aircraft_category",y = "count",p
plt.title("aircraft_category")
plt.ylabel("number of aircraft")
plt.xticks(rotation = 90)
plt.show()
```




```
In [60]: df_filtered = df[(df['aircraft_category'] != 'airplane') & (df['aircraft_ca

plt.figure(figsize=(6,4))
sns.countplot(data = df_filtered, x = "aircraft_category",dodge=False, pale
plt.title("aircraft_category")
plt.ylabel("number of aircraft")
plt.xticks(rotation = 90)
plt.show()
```




```
In [70]: # Remove 'unknown', 'unk' to exclude
df = df[df['aircraft_damage'] != 'UNK']
df = df[df['aircraft_damage'] != 'Unknown']
# Aggregating the data by aircraft category
df1= df['country'].value_counts().head(10).index
#df1 = df[df['country'] != 'Unknown']
df1 = df[df['country'].isin(df1)]
df2= df1['country'].value_counts().tail(9).index
#df = df[df['country'] != 'Unknown']
df2 = df[df['country'].isin(df2)]
```

```
In [69]: # Plot the bar plot
damage_data = df1.groupby(['country', 'aircraft_damage']).size().unstack().

# Plot grouped bar chart
damage_data.plot(kind='bar', figsize=(12, 6), colormap='viridis')
```


Out[69]: <Axes: xlabel='country'>



```
In [71]: # Plot the bar plot
damage_data = df2.groupby(['country', 'aircraft_damage']).size().unstack().

# Plot grouped bar chart
damage_data.plot(kind='bar', figsize=(12, 6), colormap='viridis')
```


Out[71]: <Axes: xlabel='country'>



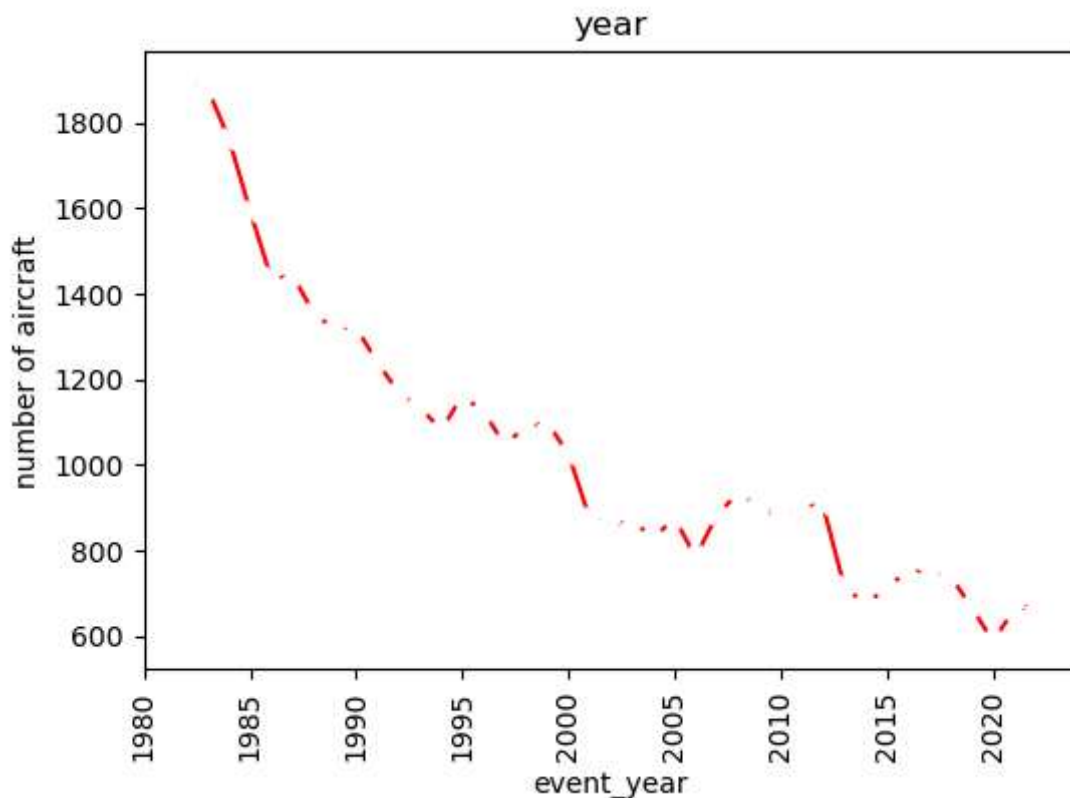
```
In [74]: # Extract the year from the event.date column for temporal analysis.  
df['event_year'] = pd.to_datetime(df['event_date']).dt.year  
df['event_year'].head()
```

```
Out[74]: 7      1982  
        10     1982  
        11     1982  
        13     1982  
        14     1982  
        Name: event_year, dtype: int32
```

```
In [78]: year = df['event_year'].value_counts().reset_index()  
year  
plt.figure(figsize=(6,4))  
sns.lineplot(data = year, x = "event_year",y = "count",palette='flare')  
plt.title("year")  
plt.ylabel("number of aircraft")  
plt.xticks(rotation = 90)  
plt.show()
```



```
In [81]: year = df['event_year'].value_counts().reset_index()
year
plt.figure(figsize=(6,4))
sns.lineplot(data = year, x = "event_year",y = "count",marker = "x", marker
plt.title("year")
plt.ylabel("number of aircraft")
plt.xticks(rotation = 90)
plt.show()
```



```
In [ ]: from wordcloud import WordCloud
# Remove rows where airport.name contains 'Unknown' or 'None'
df = df[~df['airport.name'].str.contains('Unknown|None', case=False, na=False)]

# Filter data for high-risk airports
high_risk_airports = df[df['risk.category'] == 'High Risk']['airport.name']

# Generate word cloud
wordcloud = WordCloud(width=800, height=400, background_color='white').generate_from_texts(
    high_risk_airports

# Plot
plt.figure(figsize=(12, 8))
plt.imshow(wordcloud, interpolation='lanczos')
plt.axis('off')
plt.title('High-Risk Airports', fontsize=16, fontweight='bold')
plt.show()
```


```

In [82]: df1= df['make'].value_counts().head(3).index
#df1 = df[df['make'] != 'Unknown']
df1 = df[df['make'].isin(df1)]
damage_data = df1.groupby(['make', 'broad_phase_of_flight']).size().unstack

# Plot grouped bar chart
damage_data.plot(kind='bar', figsize=(12, 6), colormap='viridis')

```

Out[82]: <Axes: xlabel='make'>



In []: